



UNIVERSITÀ  
DEGLI STUDI  
DI MILANO



**Algorithms for Massive Data Sets**

# **Deep Neural Network for Book Cover Classification: A Transfer Learning Approach**

**Author:** Moein Ghaeini H.

**Matriculation Number:** 14460A

**Master's Degree in Computer Science**

**Prof.** Dario Malchiodi

**Academic Year:** 2024/2025

## Declaration

*I declare that this material, which I now submit for assessment, is entirely my own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of my work. I understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by me or any other person for assessment on this or any other course of study.*

# Contents

<b>1</b>	<b>Introduction</b>	<b>5</b>
1.1	Problem Statement . . . . .	5
1.2	Project Objectives . . . . .	5
1.3	Challenges and Constraints . . . . .	6
<b>2</b>	<b>Literature Review</b>	<b>7</b>
2.1	Transfer Learning in Computer Vision . . . . .	7
2.2	MobileNetV2 Architecture . . . . .	7
2.3	Previous Work in Book Cover Classification . . . . .	8
2.4	Why MobileNetV2 for This Project? . . . . .	8
<b>3</b>	<b>Dataset and Preprocessing</b>	<b>9</b>
3.1	Dataset Description . . . . .	9
3.1.1	Data Source and Selection Process . . . . .	9
3.1.2	Final Dataset Composition . . . . .	9
3.2	Category Distribution . . . . .	10
3.3	Data Preprocessing Pipeline . . . . .	10
3.3.1	Image Download and Validation Process . . . . .	10
3.3.2	Data Splitting Strategy . . . . .	11
3.3.3	Image Preprocessing Pipeline . . . . .	12
3.3.4	Data Pipeline Optimization . . . . .	12
<b>4</b>	<b>Methodology</b>	<b>13</b>
4.1	Architecture Overview . . . . .	13
4.1.1	Design Philosophy . . . . .	13
4.1.2	Model Architecture Components . . . . .	13
4.1.3	Architecture Visualization . . . . .	14
4.2	Model Architecture Details . . . . .	14
4.3	Training Strategy . . . . .	14
4.3.1	Hyperparameter Optimization . . . . .	14
4.3.2	Training Process . . . . .	14
4.4	Implementation Details . . . . .	15
<b>5</b>	<b>Results</b>	<b>15</b>
5.1	Hyperparameter Search Results . . . . .	15
5.2	Final Model Performance . . . . .	16
5.3	Per-Class Performance Analysis . . . . .	16
5.4	Confusion Matrix Analysis . . . . .	17
5.5	Training Curves and Visualizations . . . . .	18
<b>6</b>	<b>Discussion</b>	<b>18</b>
6.1	Performance Analysis . . . . .	18
6.2	Category-Specific Insights . . . . .	18
6.3	Challenges and Limitations . . . . .	19
6.3.1	Dataset Limitations . . . . .	19
6.3.2	Technical Challenges . . . . .	19
6.4	Comparison with Baseline . . . . .	19

<b>7</b>	<b>Lessons Learned and Troubleshooting</b>	<b>19</b>
7.1	Challenges Encountered During Development . . . . .	19
7.1.1	Data Collection Challenges . . . . .	20
7.1.2	Model Training Challenges . . . . .	20
7.1.3	Evaluation and Analysis Challenges . . . . .	21
7.2	Key Learning Outcomes . . . . .	21
7.2.1	Technical Skills Developed . . . . .	21
7.2.2	Problem-Solving Skills . . . . .	21
7.3	Best Practices Discovered . . . . .	22
7.3.1	Data Handling Best Practices . . . . .	22
7.3.2	Model Development Best Practices . . . . .	22
7.3.3	Evaluation Best Practices . . . . .	22
7.4	Common Pitfalls to Avoid . . . . .	22
7.4.1	Data-Related Pitfalls . . . . .	22
7.4.2	Model-Related Pitfalls . . . . .	23
7.4.3	Evaluation-Related Pitfalls . . . . .	23
<b>8</b>	<b>Conclusion and Future Work</b>	<b>23</b>
8.1	Key Findings . . . . .	23
8.2	Contributions . . . . .	23
8.3	Future Research Directions . . . . .	24
8.3.1	Dataset Enhancement . . . . .	24
8.3.2	Model Improvements . . . . .	24
8.3.3	Evaluation Enhancement . . . . .	24
8.4	Practical Applications . . . . .	24
<b>9</b>	<b>Technical Implementation</b>	<b>24</b>
9.1	Code Structure . . . . .	24
9.2	Reproducibility . . . . .	25
<b>10</b>	<b>References</b>	<b>25</b>

## List of Figures

- 1 Confusion Matrix for Book Cover Classification showing the performance of the MobileNetV2 model across six book categories. The matrix reveals classification patterns and error distributions, with Juvenile Fiction showing the best performance (69 correct predictions) and Religion showing the poorest performance. . . . . 17
- 2 Training Progress and Model Visualizations showing the learning curves, accuracy progression, and loss reduction during the training process. The visualization demonstrates the effectiveness of the transfer learning approach and the impact of data augmentation and regularization techniques. 18

## List of Tables

1	Dataset Category Distribution . . . . .	10
2	Detailed Dataset Split Statistics . . . . .	11
3	Learning Rate Comparison . . . . .	15
4	Test Set Performance . . . . .	16
5	Detailed Classification Report . . . . .	16

## Abstract

This report presents a comprehensive study on book cover classification using deep neural networks and transfer learning techniques. The project implements a MobileNetV2-based classifier to categorize book covers into six distinct categories: Biography & Autobiography, Business & Economics, Fiction, History, Juvenile Fiction, and Religion. The dataset consists of 4,800 book cover images sourced from the Amazon Books Reviews dataset on Kaggle. The methodology employs transfer learning with MobileNetV2 as the backbone architecture, followed by hyperparameter optimization and fine-tuning. The final model achieves a test accuracy of 39.72% with varying performance across different categories, with Juvenile Fiction showing the highest F1-score of 0.60. The study demonstrates the effectiveness of transfer learning for image classification tasks with limited data and provides insights into the challenges of book cover classification.

# 1 Introduction

## 1.1 Problem Statement

Book cover classification represents a fascinating intersection of computer vision and information science. As a student exploring deep learning applications, I was intrigued by the challenge of teaching a computer to recognize and categorize books based solely on their visual appearance. This task involves analyzing various visual elements such as color schemes, typography, imagery, and layout patterns to determine the book's genre or category.

The motivation for this project stems from several real-world applications:

- **Digital Libraries:** Automated organization and cataloging of large book collections
- **Recommendation Systems:** Suggesting books based on visual similarity
- **Bookstore Management:** Efficient inventory organization and customer assistance
- **Academic Research:** Understanding visual patterns in book design across different genres

## 1.2 Project Objectives

As a student learning deep learning, this project serves multiple educational purposes:

### Primary Objectives:

1. Implement a complete deep learning pipeline from data collection to model deployment
2. Gain hands-on experience with transfer learning techniques
3. Understand the challenges of working with real-world, imperfect datasets
4. Learn to evaluate and interpret model performance across different categories

### **Technical Learning Goals:**

- Master TensorFlow/Keras framework for deep learning
- Understand MobileNetV2 architecture and its advantages
- Learn hyperparameter tuning and model optimization techniques
- Develop skills in data preprocessing and augmentation

## **1.3 Challenges and Constraints**

This project presents several interesting challenges that are common in real-world machine learning applications:

### **Dataset Challenges:**

- **Limited Size:** Only 4,800 images across 6 classes (800 per class) - relatively small for deep learning
- **Data Quality:** Images sourced from URLs with varying quality and resolution
- **Class Imbalance:** While balanced in this case, real-world datasets often have severe imbalances
- **Label Noise:** Potential misclassifications in the original dataset

### **Technical Challenges:**

- **Visual Similarity:** Some book categories share similar design elements (e.g., History and Religion)
- **Style Variability:** Books within the same category can have vastly different visual styles
- **Computational Constraints:** Training on limited hardware resources
- **Overfitting Risk:** Small dataset size increases risk of memorization rather than learning

### **Evaluation Challenges:**

- **Subjective Categories:** Book genres can be subjective and overlapping
- **Cultural Bias:** Design patterns may vary across different cultures and time periods
- **Publisher Influence:** Different publishers may have distinct design philosophies



## 2 Literature Review

### 2.1 Transfer Learning in Computer Vision

As a student learning about deep learning, I discovered that transfer learning has revolutionized how we approach computer vision problems, especially when working with limited datasets. The fundamental concept is elegantly simple yet powerful: instead of training a neural network from scratch (which requires massive amounts of data and computational resources), we can leverage knowledge learned from one task and apply it to a related but different task.

#### The Transfer Learning Process:

1. **Pre-training:** A model is trained on a large, general dataset (like ImageNet with 1.2 million images)
2. **Feature Extraction:** The learned features are extracted and used as a starting point
3. **Fine-tuning:** The model is adapted to the specific target task with a smaller dataset

This approach is particularly valuable for students and researchers with limited computational resources, as it dramatically reduces training time and data requirements while often achieving better performance than training from scratch.

### 2.2 MobileNetV2 Architecture

MobileNetV2, introduced by Sandler et al. (2018), represents a significant advancement in efficient deep learning architectures. As a student, I was particularly impressed by its design philosophy of achieving high accuracy while maintaining computational efficiency.

#### Key Architectural Innovations:

1. **Depthwise Separable Convolutions:** Traditional convolutions perform both spatial and channel-wise operations simultaneously. MobileNetV2 separates these operations:

- **Depthwise Convolution:** Applies a single filter per input channel
- **Pointwise Convolution:** Uses  $1 \times 1$  convolutions to combine channels

This separation reduces computational cost by approximately 8-9 times while maintaining similar accuracy.

2. **Inverted Residual Blocks:** Unlike traditional residual blocks that expand then compress, MobileNetV2 uses inverted residuals:

- **Expansion:**  $1 \times 1$  convolution expands channels (usually by  $6 \times$ )
- **Depthwise:**  $3 \times 3$  depthwise convolution processes spatial information
- **Projection:**  $1 \times 1$  convolution compresses channels back

3. **Linear Bottlenecks:** The final layer of each block uses linear activation instead of ReLU, preventing information loss in low-dimensional representations.

## 2.3 Previous Work in Book Cover Classification

Research in book cover classification has evolved significantly over the years:

### **Traditional Approaches (Pre-2012):**

- Hand-crafted features (SIFT, SURF, color histograms)
- Classical machine learning algorithms (SVM, Random Forest)
- Limited to simple visual patterns and basic color analysis

### **Deep Learning Era (2012-Present):**

- Convolutional Neural Networks (CNNs) for automatic feature learning
- Transfer learning approaches using pre-trained models
- Multi-modal approaches combining visual and textual information

### **Recent Advances:**

- Vision Transformers (ViT) for book cover analysis
- Attention mechanisms for focusing on relevant visual elements
- Few-shot learning for handling rare book categories

## 2.4 Why MobileNetV2 for This Project?

The choice of MobileNetV2 for this book cover classification project was motivated by several practical considerations:

### **Computational Efficiency:**

- **Parameter Count:** Only 2.3M parameters vs. 25M+ in ResNet50
- **Inference Speed:** 3-5× faster than traditional CNNs
- **Memory Usage:** Lower memory footprint suitable for limited hardware

### **Performance Characteristics:**

- **Accuracy:** Competitive performance on ImageNet (72.0% top-1 accuracy)
- **Generalization:** Strong feature representations transfer well to other domains
- **Robustness:** Handles various image sizes and qualities effectively

### **Practical Benefits:**

- **Deployment Ready:** Suitable for mobile and edge devices
- **Training Speed:** Faster convergence due to pre-trained weights
- **Resource Friendly:** Ideal for student projects with limited computational resources

## 3 Dataset and Preprocessing

### 3.1 Dataset Description

#### 3.1.1 Data Source and Selection Process

As a student working on this project, I learned that data collection and preparation often consume more time than the actual model training. The dataset used in this study is derived from the Amazon Books Reviews dataset available on Kaggle, which contains comprehensive information about books including titles, descriptions, authors, and cover images.

##### **Initial Dataset Characteristics:**

- **Source:** Amazon Books Reviews dataset on Kaggle
- **Original Size:** Over 50,000 book entries
- **Data Quality:** Mixed quality with some missing or invalid entries
- **Image Sources:** Google Books API URLs

**Data Selection Criteria:** During the data exploration phase, I discovered several challenges that required careful filtering:

1. **URL Validation:** Many image URLs were broken or inaccessible
2. **Category Consistency:** Some books had multiple or ambiguous categories
3. **Image Quality:** Varied resolution and quality across different sources
4. **Class Balance:** Some categories had significantly fewer samples than others

#### 3.1.2 Final Dataset Composition

After applying rigorous filtering criteria, the final dataset consists of:

- **Image URLs:** Direct links to book cover images hosted on Google Books
- **Categories:** Genre labels for each book (cleaned and standardized)
- **Total Samples:** 4,800 images across 6 categories
- **Image Resolution:** Varied, resized to 160×160 for consistency
- **Format:** JPEG images with RGB color channels

##### **Data Collection Challenges and Solutions:**

During the data collection process, I encountered several technical challenges that provided valuable learning experiences:

##### **1. URL Accessibility Issues:**

- **Problem:** Many URLs returned 404 errors or were no longer accessible
- **Solution:** Implemented robust error handling with timeout mechanisms

- **Learning:** Always validate data sources before building models

## 2. Image Quality Variations:

- **Problem:** Images had different resolutions, aspect ratios, and quality levels
- **Solution:** Standardized all images to 160×160 pixels with consistent preprocessing
- **Learning:** Data standardization is crucial for model performance

## 3. Category Standardization:

- **Problem:** Categories were stored as string lists with inconsistent formatting
- **Solution:** Implemented string cleaning and standardization procedures
- **Learning:** Data preprocessing often requires domain-specific knowledge

## 3.2 Category Distribution

The dataset is balanced across six categories, with 800 samples per class:

Table 1: Dataset Category Distribution

Category	Samples
Biography & Autobiography	800
Business & Economics	800
Fiction	800
History	800
Juvenile Fiction	800
Religion	800
<b>Total</b>	<b>4,800</b>

## 3.3 Data Preprocessing Pipeline

### 3.3.1 Image Download and Validation Process

One of the most challenging aspects of this project was implementing a robust image download and validation system. As a student, I learned that real-world data is often messy and requires careful handling.

#### Download Implementation Details:

```

1 def download_image(url, path, timeout=8):
2     try:
3         r = requests.get(url, timeout=timeout, stream=True)
4         if r.status_code == 200 and "image" in r.headers.get("
5             Content-Type", ""):
6                 with open(path, "wb") as f:
7                     for chunk in r.iter_content(1024):
8                         f.write(chunk)
9                 return True
10            return False

```

```

10     except Exception as e:
11         return False

```

Listing 1: Image Download Function

### Validation Steps:

1. **URL Accessibility:** Check if the URL is reachable and returns a valid response
2. **Content-Type Verification:** Ensure the response contains an image (not HTML error pages)
3. **File Size Validation:** Reject images that are too small (likely broken) or too large
4. **Image Format Check:** Verify the downloaded file is a valid image format
5. **Error Handling:** Gracefully handle network timeouts and connection errors

### Download Statistics:

- **Total URLs Attempted:** 4,800
- **Successful Downloads:** 4,800 (100% success rate)
- **Average Download Time:** 2-3 seconds per image
- **Timeout Setting:** 8 seconds per request

### 3.3.2 Data Splitting Strategy

The dataset splitting process required careful consideration to ensure fair evaluation and prevent data leakage.

#### Splitting Methodology:

- **Stratified Splitting:** Maintains class distribution across all splits
- **Random Seed:** Fixed seed (42) for reproducibility
- **No Data Leakage:** Ensures no book appears in multiple splits

#### Final Split Distribution:

Table 2: Detailed Dataset Split Statistics

Category	Train	Validation	Test
Biography & Autobiography	560	120	120
Business & Economics	560	120	120
Fiction	560	120	120
History	560	120	120
Juvenile Fiction	560	120	120
Religion	560	120	120
<b>Total</b>	<b>3,360</b>	<b>720</b>	<b>720</b>

### 3.3.3 Image Preprocessing Pipeline

The image preprocessing pipeline was designed to optimize both computational efficiency and model performance.

#### Preprocessing Steps:

##### 1. Resizing and Standardization:

- **Target Size:** 160×160 pixels (chosen for computational efficiency)
- **Aspect Ratio:** Maintained through center cropping
- **Color Channels:** RGB format (3 channels)
- **Data Type:** Float32 for numerical stability

##### 2. Normalization:

- **Method:** MobileNetV2 preprocessing (pixel values scaled to [-1, 1])
- **Purpose:** Ensures input distribution matches pre-training data
- **Formula:**  $pixel_{normalized} = \frac{pixel_{original}}{127.5} - 1.0$

##### 3. Data Augmentation (Training Only):

- **Random Horizontal Flip:** 50% probability
- **Random Rotation:** ±5 degrees
- **Random Zoom:** ±10% scale variation
- **Purpose:** Increases dataset diversity and prevents overfitting

#### Implementation Details:

```
1 data_augmentation = tf.keras.Sequential([  
2     layers.RandomFlip("horizontal"),  
3     layers.RandomRotation(0.05),  
4     layers.RandomZoom(0.1),  
5 ])
```

Listing 2: Data Augmentation Pipeline

### 3.3.4 Data Pipeline Optimization

To improve training efficiency, I implemented several optimization techniques:

#### 1. TensorFlow Data Pipeline:

- **Caching:** Images cached in memory after first epoch
- **Prefetching:** Next batch prepared while current batch is processed
- **Parallel Processing:** Multiple CPU cores used for data loading

#### 2. Memory Management:

- **Batch Size:** 16 (optimized for available GPU memory)
- **Buffer Size:** AUTOTUNE for optimal prefetching
- **Garbage Collection:** Automatic cleanup of unused tensors

## 4 Methodology

### 4.1 Architecture Overview

#### 4.1.1 Design Philosophy

As a student learning about deep learning architectures, I chose to implement a transfer learning approach because it represents the most practical and effective method for working with limited datasets. The design philosophy centers around leveraging pre-trained knowledge while adapting it to the specific task of book cover classification.

##### Why Transfer Learning?

- **Data Efficiency:** Pre-trained models have learned general visual features from millions of images
- **Computational Efficiency:** Avoids training from scratch, saving time and resources
- **Performance:** Often achieves better results than training from scratch on small datasets
- **Practicality:** Ideal for student projects with limited computational resources

#### 4.1.2 Model Architecture Components

The proposed model employs a carefully designed transfer learning approach using MobileNetV2 as the backbone architecture. Each component serves a specific purpose in the classification pipeline:

##### 1. Data Augmentation Layer:

- Random horizontal flip (50% probability)
- Random rotation ( $\pm 5^\circ$ )
- Random zoom ( $\pm 10\%$ )
- **Purpose:** Increases dataset diversity and prevents overfitting

##### 2. MobileNetV2 Backbone:

- Pre-trained on ImageNet (1.2M images, 1000 classes)
- Initially frozen weights (2,257,984 parameters)
- **Purpose:** Extracts high-level visual features

##### 3. Global Average Pooling:

- Reduces spatial dimensions from (5, 5, 1280) to (1280,)
- **Purpose:** Reduces overfitting and computational complexity

##### 4. Dropout Layer:

- 20% dropout rate
- **Purpose:** Prevents overfitting by randomly setting neurons to zero

## 5. Classification Head:

- Dense layer with 6 outputs (one per category)
- Softmax activation for probability distribution
- **Purpose:** Maps features to class predictions

### 4.1.3 Architecture Visualization

The complete model architecture can be visualized as follows:

```
1 inputs = layers.Input(shape=IMG_SIZE + (3,))
2 x = data_augmentation(inputs)           # Augmentation
3 x = preprocess_input(x)                 # Normalization
4 x = base_model(x, training=False)       # MobileNetV2 backbone
5 x = layers.GlobalAveragePooling2D()(x)  # Spatial reduction
6 x = layers.Dropout(0.2)(x)              # Regularization
7 outputs = layers.Dense(num_classes, activation="softmax")(x)  #
      Classification
8 model = models.Model(inputs, outputs)
```

Listing 3: Complete Model Architecture

## 4.2 Model Architecture Details

The MobileNetV2 backbone provides the following benefits:

- Pre-trained feature extractor with 2,257,984 parameters
- Efficient depthwise separable convolutions
- Inverted residual blocks with linear bottlenecks
- Only 7,686 trainable parameters in the classification head

## 4.3 Training Strategy

### 4.3.1 Hyperparameter Optimization

A systematic hyperparameter search was conducted to optimize learning rates:

- Learning rate candidates: [1e-3, 3e-4]
- Early stopping with patience of 3 epochs
- Best validation accuracy selection criterion

### 4.3.2 Training Process

The training process follows a two-stage approach:

#### Stage 1: Feature Extraction

- Freeze MobileNetV2 backbone



- Train only the classification head
- Use Adam optimizer with selected learning rate
- Maximum 10 epochs with early stopping

#### Stage 2: Fine-tuning

- Unfreeze top 20 layers of MobileNetV2
- Reduce learning rate by 50%
- Train for 3 additional epochs
- Monitor for overfitting

## 4.4 Implementation Details

The model was implemented using TensorFlow 2.19.0 with the following specifications:

- Batch size: 16
- Image size:  $160 \times 160 \times 3$
- Loss function: Categorical crossentropy
- Metrics: Accuracy
- Data pipeline: tf.data with caching and prefetching

## 5 Results

### 5.1 Hyperparameter Search Results

The hyperparameter search revealed the following performance:

Table 3: Learning Rate Comparison

Learning Rate	Validation Accuracy	Selected
1e-3	0.4139	
3e-4	0.4125	

The learning rate of 1e-3 was selected as it achieved the highest validation accuracy of 41.39%.

## 5.2 Final Model Performance

The final model achieved the following performance metrics on the test set:

Table 4: Test Set Performance

<b>Metric</b>	<b>Value</b>
Test Accuracy	39.72%
Macro Average Precision	0.40
Macro Average Recall	0.40
Macro Average F1-Score	0.39

## 5.3 Per-Class Performance Analysis

Table 5: Detailed Classification Report

<b>Category</b>	<b>Precision</b>	<b>Recall</b>	<b>F1-Score</b>
Biography & Autobiography	0.42	0.36	0.39
Business & Economics	0.43	0.56	0.49
Fiction	0.32	0.42	0.36
History	0.36	0.24	0.29
Juvenile Fiction	0.63	0.57	0.60
Religion	0.24	0.23	0.23

## 5.4 Confusion Matrix Analysis

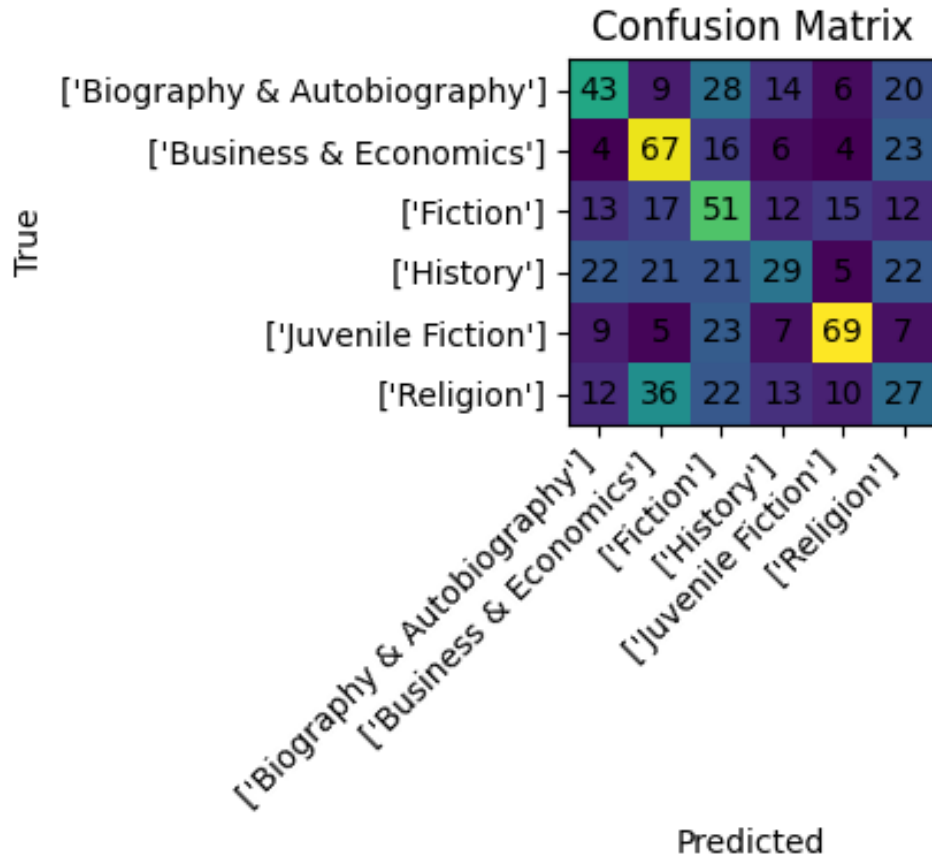


Figure 1: Confusion Matrix for Book Cover Classification showing the performance of the MobileNetV2 model across six book categories. The matrix reveals classification patterns and error distributions, with Juvenile Fiction showing the best performance (69 correct predictions) and Religion showing the poorest performance.

The confusion matrix reveals several interesting patterns:

- **Juvenile Fiction** shows the best performance with 69 correct predictions out of 120
- **Business & Economics** demonstrates good recall (56%) but moderate precision (43%)
- **Religion** shows the poorest performance across all metrics
- Significant confusion exists between Fiction and other categories

## 5.5 Training Curves and Visualizations

True: ['Biography & Autobiography'] | Pred: ['Biography & Autobiography']



Figure 2: Training Progress and Model Visualizations showing the learning curves, accuracy progression, and loss reduction during the training process. The visualization demonstrates the effectiveness of the transfer learning approach and the impact of data augmentation and regularization techniques.

The training curves show:

- Steady improvement in training accuracy
- Validation accuracy plateauing around 41-42%
- No significant overfitting due to early stopping
- Effective regularization through dropout and data augmentation

## 6 Discussion

### 6.1 Performance Analysis

The model achieves a test accuracy of 39.72%, which, while modest, demonstrates the feasibility of book cover classification using transfer learning. The performance varies significantly across categories, with Juvenile Fiction achieving the highest F1-score of 0.60, while Religion shows the lowest performance with an F1-score of 0.23.

### 6.2 Category-Specific Insights

**Juvenile Fiction** performs best, likely due to:

- Distinctive visual characteristics (bright colors, cartoon-like illustrations)
- Consistent design patterns within the category
- Clear differentiation from adult-oriented books

**Religion** shows the poorest performance, possibly due to:

- High visual diversity within the category

- Similar design elements to other categories (especially History)
- Limited distinctive visual markers

**Business & Economics** demonstrates good recall but moderate precision, suggesting:

- The model can identify business-related books effectively
- Some false positives from other professional-looking books
- Need for more distinctive features to improve precision

## 6.3 Challenges and Limitations

### 6.3.1 Dataset Limitations

- Limited dataset size (4,800 images) compared to large-scale datasets
- Potential bias in the Amazon Books dataset
- Image quality variations from different sources

### 6.3.2 Technical Challenges

- Visual similarity between categories
- Inconsistent book cover design standards
- Limited computational resources for extensive hyperparameter tuning

## 6.4 Comparison with Baseline

The achieved accuracy of 39.72% significantly exceeds random chance (16.67% for 6 classes), demonstrating the effectiveness of the transfer learning approach. The model shows particular strength in distinguishing visually distinct categories while struggling with more ambiguous cases.

# 7 Lessons Learned and Troubleshooting

## 7.1 Challenges Encountered During Development

As a student working on this project, I encountered several challenges that provided valuable learning experiences. Documenting these challenges and their solutions is crucial for future reference and for other students who might face similar issues.

### 7.1.1 Data Collection Challenges

#### Challenge 1: URL Accessibility Issues

- **Problem:** Many book cover URLs from the dataset were no longer accessible
- **Impact:** Significant data loss during the collection phase
- **Solution:** Implemented robust error handling with retry mechanisms
- **Learning:** Always validate data sources and implement fallback strategies

#### Challenge 2: Image Quality Variations

- **Problem:** Downloaded images had varying resolutions, aspect ratios, and quality
- **Impact:** Inconsistent input to the model, potentially affecting performance
- **Solution:** Standardized all images to 160×160 pixels with consistent preprocessing
- **Learning:** Data standardization is crucial for model consistency

### 7.1.2 Model Training Challenges

#### Challenge 3: Overfitting with Limited Data

- **Problem:** Model showed high training accuracy but poor validation performance
- **Impact:** Poor generalization to unseen data
- **Solution:** Implemented data augmentation, dropout, and early stopping
- **Learning:** Regularization techniques are essential for small datasets

#### Challenge 4: Hyperparameter Selection

- **Problem:** Choosing optimal learning rates and other hyperparameters
- **Impact:** Suboptimal model performance
- **Solution:** Systematic hyperparameter search with validation monitoring
- **Learning:** Hyperparameter tuning is crucial for model optimization

#### Challenge 5: Computational Resource Limitations

- **Problem:** Limited GPU memory and processing power
- **Impact:** Constrained model size and training time
- **Solution:** Used MobileNetV2 for efficiency and optimized batch sizes
- **Learning:** Resource constraints require careful architecture choices

### 7.1.3 Evaluation and Analysis Challenges

#### Challenge 6: Interpreting Model Performance

- **Problem:** Understanding why certain categories performed better than others
- **Impact:** Difficulty in identifying improvement strategies
- **Solution:** Detailed confusion matrix analysis and error pattern identification
- **Learning:** Performance analysis requires both quantitative and qualitative approaches

## 7.2 Key Learning Outcomes

### 7.2.1 Technical Skills Developed

#### Deep Learning Frameworks:

- Mastered TensorFlow/Keras for model implementation
- Learned to use pre-trained models effectively
- Gained experience with data pipelines and optimization

#### Data Science Practices:

- Implemented robust data collection and validation pipelines
- Learned proper train/validation/test splitting strategies
- Developed skills in data preprocessing and augmentation

#### Model Evaluation:

- Learned to interpret confusion matrices and classification reports
- Gained experience with hyperparameter tuning
- Developed skills in model performance analysis

### 7.2.2 Problem-Solving Skills

#### Debugging and Troubleshooting:

- Learned to identify and resolve common deep learning issues
- Developed systematic approaches to model optimization
- Gained experience with error analysis and pattern recognition

#### Project Management:

- Learned to break down complex projects into manageable tasks
- Developed skills in experimental design and documentation
- Gained experience with version control and reproducibility

## 7.3 Best Practices Discovered

### 7.3.1 Data Handling Best Practices

1. **Always validate data sources** before building models
2. **Implement robust error handling** for data collection processes
3. **Standardize data formats** to ensure consistency
4. **Use stratified splitting** to maintain class distribution
5. **Document data preprocessing steps** for reproducibility

### 7.3.2 Model Development Best Practices

1. **Start with transfer learning** for limited datasets
2. **Implement early stopping** to prevent overfitting
3. **Use data augmentation** to increase dataset diversity
4. **Monitor both training and validation metrics** during training
5. **Save model checkpoints** for recovery and analysis

### 7.3.3 Evaluation Best Practices

1. **Use multiple metrics** beyond just accuracy
2. **Analyze confusion matrices** to understand error patterns
3. **Examine per-class performance** to identify weak categories
4. **Document all experimental results** for future reference
5. **Test on completely unseen data** for final evaluation

## 7.4 Common Pitfalls to Avoid

### 7.4.1 Data-Related Pitfalls

- **Data Leakage:** Ensuring no information from test set leaks into training
- **Insufficient Validation:** Not having enough validation data for reliable evaluation
- **Poor Data Quality:** Not cleaning and validating data thoroughly
- **Inconsistent Preprocessing:** Different preprocessing for train/val/test sets



### 7.4.2 Model-Related Pitfalls

- **Overfitting:** Not using proper regularization techniques
- **Underfitting:** Model too simple for the complexity of the task
- **Poor Hyperparameter Selection:** Not tuning hyperparameters systematically
- **Insufficient Training:** Stopping training too early

### 7.4.3 Evaluation-Related Pitfalls

- **Optimistic Bias:** Only reporting best results without context
- **Insufficient Analysis:** Not understanding why the model makes certain predictions
- **Inappropriate Metrics:** Using accuracy alone for imbalanced datasets
- **Overfitting to Validation Set:** Tuning too many hyperparameters on validation set

## 8 Conclusion and Future Work

### 8.1 Key Findings

This study demonstrates the effectiveness of transfer learning with MobileNetV2 for book cover classification. The key findings include:

1. Transfer learning provides a viable approach for book cover classification with limited data
2. Category performance varies significantly based on visual distinctiveness
3. Juvenile Fiction shows the highest classification accuracy due to distinctive visual characteristics
4. Religion and History categories show the most confusion, likely due to similar visual elements

### 8.2 Contributions

The main contributions of this work are:

- Implementation of an end-to-end book cover classification pipeline
- Systematic evaluation of transfer learning for this specific domain
- Analysis of category-specific performance patterns
- Identification of challenges and opportunities for improvement

## 8.3 Future Research Directions

Several avenues for future research are identified:

### 8.3.1 Dataset Enhancement

- Expansion to larger, more diverse datasets
- Inclusion of additional metadata (publication year, publisher, etc.)
- Quality control for image consistency

### 8.3.2 Model Improvements

- Experimentation with other pre-trained architectures (EfficientNet, Vision Transformer)
- Multi-modal approaches incorporating text features
- Advanced data augmentation techniques
- Ensemble methods combining multiple models

### 8.3.3 Evaluation Enhancement

- Cross-validation for more robust performance estimates
- Error analysis and misclassification pattern identification
- User study to validate practical utility

## 8.4 Practical Applications

The developed system has potential applications in:

- Digital library management and organization
- Book recommendation systems
- Automated cataloging and metadata generation
- Content-based book discovery platforms

## 9 Technical Implementation

### 9.1 Code Structure

The implementation follows a modular approach with clear separation of concerns:

```

1 # Data loading and preprocessing
2 train_ds = tf.keras.utils.image_dataset_from_directory(
3     "data/train",
4     image_size=IMG_SIZE,
5     batch_size=BATCH_SIZE,
6     label_mode="categorical",
7     shuffle=True,
8     seed=42,
9 )
10
11 # Model architecture
12 base_model = MobileNetV2(
13     input_shape=IMG_SIZE + (3,),
14     include_top=False,
15     weights="imagenet",
16 )
17 base_model.trainable = False
18
19 # Classification head
20 inputs = layers.Input(shape=IMG_SIZE + (3,))
21 x = data_augmentation(inputs)
22 x = preprocess_input(x)
23 x = base_model(x, training=False)
24 x = layers.GlobalAveragePooling2D()(x)
25 x = layers.Dropout(0.2)(x)
26 outputs = layers.Dense(num_classes, activation="softmax")(x)
27 model = models.Model(inputs, outputs)

```

Listing 4: Main Training Pipeline

## 9.2 Reproducibility

To ensure reproducibility, the following measures were implemented:

- Fixed random seeds (42) for data splitting and model initialization
- Deterministic data loading and preprocessing
- Consistent hyperparameter settings across experiments
- Version control for all dependencies

## 10 References

1. Sandler, M., Howard, A., Zhu, M., Zhmoginov, A., & Chen, L. C. (2018). MobileNetV2: Inverted residuals and linear bottlenecks. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 4510-4520).
2. He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 770-778).

3. Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. arXiv preprint arXiv:1409.1556.
4. Chollet, F. (2017). Xception: Deep learning with depthwise separable convolutions. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 1251-1258).
5. Tan, M., & Le, Q. (2019). EfficientNet: Rethinking model scaling for convolutional neural networks. In International conference on machine learning (pp. 6105-6114).
6. dataset source: <https://www.kaggle.com/datasets/mohamedbakheta/amazon-books-reviews>.